

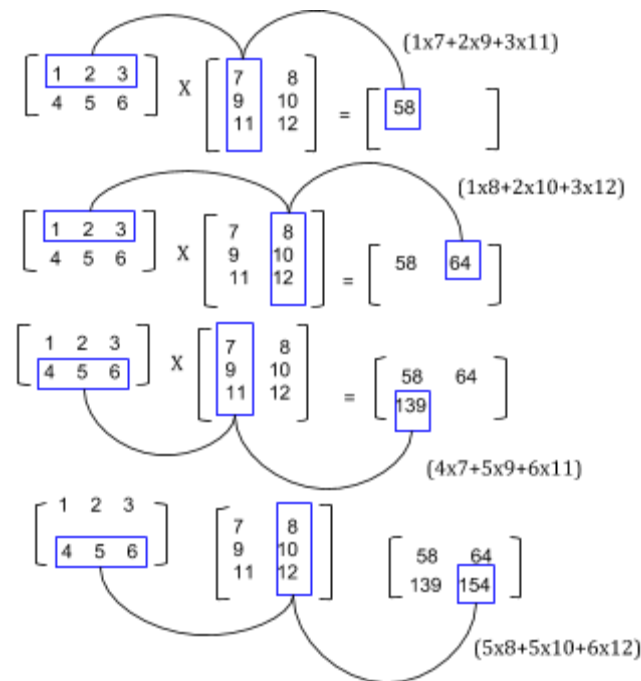
If you like my post - Do follow me on this blog -

[Click here to Follow my blog](#)

Matrix Multiplication

Using MapReduce Programming

In **mathematics**, **matrix multiplication** or the **matrix product** is a binary operation that produces a matrix from two matrices. The definition is motivated by linear equations and linear transformations on vectors, which have numerous applications in applied mathematics, physics, and engineering. In more detail, if **A** is an $n \times m$ matrix and **B** is an $m \times p$ matrix, their matrix product **AB** is an $n \times p$ matrix, in which the m entries across a row of **A** are multiplied with the m entries down a column of **B** and summed to produce an entry of **AB**. When two linear transformations are represented by matrices, then the matrix product represents the composition of the two transformations.



Algorithm for Map Function.

- for each element m_{ij} of M do
produce (key,value) pairs as $((i,k), (M,j,m_{ij}))$, for $k=1,2,3,\dots$
upto the number of columns of N
- for each element n_{jk} of N do
produce (key,value) pairs as $((i,k), (N,j,n_{jk}))$, for $i = 1,2,3,\dots$
Upto the number of rows of M .
- return Set of (key,value) pairs that each key (i,k) , has list with values (M,j,m_{ij}) and (N,j,n_{jk}) for all possible values of j .

Algorithm for Reduce Function.

for each key (i,k) do

```
sort values begin with M by j in listM
sort values begin with N by j in listN
multiply mij and njk for jth value of each list
sum up mij x njk return (i,k),  $\sum_{j=1} m_{ij} \times n_{jk}$ 
```

Step 1. Download the hadoop jar files with these links.

Download Hadoop Common Jar files: <https://goo.gl/G4MyHp>

```
$ wget https://goo.gl/G4MyHp -O hadoop-common-2.2.0.jar
```

Download Hadoop Mapreduce Jar File: <https://goo.gl/KT8yfB>

```
$ wget https://goo.gl/KT8yfB -O
```

```
hadoop-mapreduce-client-core-2.7.1.jar
```

Step 2. Creating Mapper file for Matrix Multiplication.

```
package www.ehadoopinfo.com;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

import java.io.IOException;

public class Map
    extends org.apache.hadoop.mapreduce.Mapper<LongWritable, Text, Text, Text> {
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        Configuration conf = context.getConfiguration();
        int m = Integer.parseInt(conf.get("m"));
        int p = Integer.parseInt(conf.get("p"));
        String line = value.toString();
        // (M, i, j, Mij);
        String[] indicesAndValue = line.split(",");
        Text outputKey = new Text();
        Text outputValue = new Text();
        if (indicesAndValue[0].equals("M")) {
            for (int k = 0; k < p; k++) {
                outputKey.set(indicesAndValue[1] + "," + k);
                // outputKey.set(i,k);
                outputValue.set(indicesAndValue[0] + "," +
indicesAndValue[2]
                                + "," + indicesAndValue[3]);
                // outputValue.set(M,j,Mij);
                context.write(outputKey, outputValue);
            }
        } else {
            // (N, j, k, Njk);
            for (int i = 0; i < m; i++) {
                outputKey.set(i + "," + indicesAndValue[2]);
```

```
                outputValue.set("N," + indicesAndValue[1] + "," +  
                                + indicesAndValue[3]);  
                context.write(outputKey, outputValue);  
            }  
        }  
    }  
}
```

Step 3. Creating Reducer.java file for Matrix Multiplication.

```
package www.ehadoopinfo.com;  
  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Reducer;  
  
import java.io.IOException;  
import java.util.HashMap;  
  
public class Reduce  
    extends org.apache.hadoop.mapreduce.Reducer<Text, Text, Text, Text> {  
    @Override  
    public void reduce(Text key, Iterable<Text> values, Context context)  
        throws IOException, InterruptedException {  
        String[] value;  
        //key=(i,k),  
        //Values = [(M/N,j,V/W),..]  
        HashMap<Integer, Float> hashA = new HashMap<Integer, Float>();  
        HashMap<Integer, Float> hashB = new HashMap<Integer, Float>();  
        for (Text val : values) {  
            value = val.toString().split(",");  
            if (value[0].equals("M")) {  
                hashA.put(Integer.parseInt(value[1]),  
Float.parseFloat(value[2]));  
            } else {  
                hashB.put(Integer.parseInt(value[1]),  
Float.parseFloat(value[2]));  
            }  
        }  
        int n = Integer.parseInt(context.getConfiguration().get("n"));  
        float result = 0.0f;  
        float m_ij;  
        float n_jk;  
        for (int j = 0; j < n; j++) {  
            m_ij = hashA.containsKey(j) ? hashA.get(j) : 0.0f;  
            n_jk = hashB.containsKey(j) ? hashB.get(j) : 0.0f;  
            result += m_ij * n_jk;  
        }  
        if (result != 0.0f) {  
            context.write(null,  
                new Text(key.toString() + "," +  
Float.toString(result)));  
        }  
    }  
}
```

```
    }  
}
```

Step 4. Creating MatrixMultiply.java file for

```
package www.ehadoopinfo.com;  
  
import org.apache.hadoop.conf.*;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.*;  
import org.apache.hadoop.mapreduce.*;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;  
  
public class MatrixMultiply {  
  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.err.println("Usage: MatrixMultiply <in_dir> <out_dir>");  
            System.exit(2);  
        }  
        Configuration conf = new Configuration();  
        // M is an m-by-n matrix; N is an n-by-p matrix.  
        conf.set("m", "2");  
        conf.set("n", "2");  
        conf.set("p", "2");  
        @SuppressWarnings("deprecation")  
        Job job = new Job(conf, "MatrixMultiply");  
        job.setJarByClass(MatrixMultiply.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(Text.class);  
  
        job.setMapperClass(Map.class);  
        job.setReducerClass(Reduce.class);  
  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.waitForCompletion(true);  
    }  
}
```

Step 5. Compiling the program in particular folder named as operation/

```
$ javac -cp  
hadoop-common-2.2.0.jar:hadoop-mapreduce-client-core-2.7.1.jar:operation/:. -d  
operation/ Map.java
```

```
$ javac -cp
hadoop-common-2.2.0.jar:hadoop-mapreduce-client-core-2.7.1.jar:operation/:. -d
operation/ Reduce.java
$ javac -cp
hadoop-common-2.2.0.jar:hadoop-mapreduce-client-core-2.7.1.jar:operation/:. -d
operation/ MatrixMultiply.java
```

Step 6. Let's retrieve the directory after compilation.

```
$ ls -R operation/
operation/:
www
operation/www:
ehadoopinfo
operation/www/ehadoopinfo:
com
operation/www/ehadoopinfo/com:
Map.class  MatrixMultiply.class  Reduce.class
```

Step 7. Creating Jar file for the Matrix Multiplication.

```
$ jar -cvf MatrixMultiply.jar -C operation/ .
added manifest
adding: www/(in = 0) (out= 0)(stored 0%)
adding: www/ehadoopinfo/(in = 0) (out= 0)(stored 0%)
adding: www/ehadoopinfo/com/(in = 0) (out= 0)(stored 0%)
adding: www/ehadoopinfo/com/Reduce.class(in = 2919) (out= 1271)(deflated 56%)
adding: www/ehadoopinfo/com/MatrixMultiply.class(in = 1815) (out= 932)(deflated 48%)
adding: www/ehadoopinfo/com/Map.class(in = 2353) (out= 993)(deflated 57%)
```

Step 8. Uploading the M, N file which contains the matrix multiplication data to HDFS.

```
$ cat M
M,0,0,1
M,0,1,2
M,1,0,3
M,1,1,4
$ cat N
N,0,0,5
N,0,1,6
N,1,0,7
N,1,1,8
$ hadoop fs -mkdir Matrix/
$ hadoop fs -copyFromLocal M Matrix/
$ hadoop fs -copyFromLocal N Matrix/
```

Step 9. Executing the jar file using hadoop command and thus how fetching record from HDFS and storing output in HDFS.

```
$ hadoop jar MatrixMultiply.jar www.ehadoopinfo.com.MatrixMultiply Matrix/* result/
WARNING: Use "yarn jar" to launch YARN applications.
```

```
17/10/09 14:31:22 INFO impl.TimelineClientImpl: Timeline service address:
http://sandbox.hortonworks.com:8188/ws/v1/timeline/
17/10/09 14:31:23 INFO client.RMPProxy: Connecting to ResourceManager at
sandbox.hortonworks.com/10.0.2.15:8050
17/10/09 14:31:23 WARN mapreduce.JobResourceUploader: Hadoop command-line option
parsing not performed. Implement the Tool interface and execute your application with
ToolRunner to remedy this.
17/10/09 14:31:24 INFO input.FileInputFormat: Total input paths to process : 2
17/10/09 14:31:24 INFO mapreduce.JobSubmitter: number of splits:2
17/10/09 14:31:24 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1507555978175_0006
17/10/09 14:31:25 INFO impl.YarnClientImpl: Submitted application
application_1507555978175_0006
17/10/09 14:31:25 INFO mapreduce.Job: The url to track the job:
http://sandbox.hortonworks.com:8088/proxy/application_1507555978175_0006/
17/10/09 14:31:25 INFO mapreduce.Job: Running job: job_1507555978175_0006
17/10/09 14:31:35 INFO mapreduce.Job: Job job_1507555978175_0006 running in uber mode
: false
17/10/09 14:31:35 INFO mapreduce.Job: map 0% reduce 0%
17/10/09 14:31:45 INFO mapreduce.Job: map 100% reduce 0%
17/10/09 14:31:53 INFO mapreduce.Job: map 100% reduce 100%
17/10/09 14:31:54 INFO mapreduce.Job: Job job_1507555978175_0006 completed
successfully
17/10/09 14:31:55 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=198
        FILE: Number of bytes written=386063
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=302
        HDFS: Number of bytes written=36
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=15088
        Total time spent by all reduces in occupied slots (ms)=6188
        Total time spent by all map tasks (ms)=15088
        Total time spent by all reduce tasks (ms)=6188
        Total vcore-seconds taken by all map tasks=15088
        Total vcore-seconds taken by all reduce tasks=6188
        Total megabyte-seconds taken by all map tasks=3772000
        Total megabyte-seconds taken by all reduce tasks=1547000
    Map-Reduce Framework
        Map input records=8
        Map output records=16
```

```
Map output bytes=160
Map output materialized bytes=204
Input split bytes=238
Combine input records=0
Combine output records=0
Reduce input groups=4
Reduce shuffle bytes=204
Reduce input records=16
Reduce output records=4
Spilled Records=32
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=196
CPU time spent (ms)=2720
Physical memory (bytes) snapshot=536309760
Virtual memory (bytes) snapshot=2506076160
Total committed heap usage (bytes)=360185856

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=64

File Output Format Counters
  Bytes Written=36
```

Step 10. Getting Output from part-r-00000 that was generated after the execution of the hadoop command.

```
$ hadoop fs -cat result/part-r-00000
0,0,19.0
0,1,22.0
1,0,43.0
1,1,50.0
```

In some contents I have reduced the font size to give a better code readability of this post.
Do comment me what you think about this post. If you find any errors or any suggestions or if like my post I would be glad to hear from you.

Reference:

"Matrix Multiplication with MapReduce." Lendapp, 17 May 2016,
lendap.wordpress.com/2015/02/16/matrix-multiplication-with-mapreduce/.

The link to this article is as follows:

1. Google Drive: [Blog MR5. Matrix Multiplication using MapReduce Programming in Java.](#)
2. Blog Post: [Post - MR5. Matrix Multiplication using MapReduce Programming in Java.](#)
3. Follow My Blog: [Follow Me Here.](#)

4. QR Code:

